

# Detecting Harmful Medical Advice by Analyzing the Characteristics of Retweeters

Matthew Leon

Vanderbilt University

Nashville, Tennessee, USA

matthew.leon@vanderbilt.edu

**Abstract**—I study the ability of a model to discern, for tweets about the COVID-19 crisis and based on the characteristics of users who retweet it, whether or not a given article or tweet provides beneficial medical advice or could lead to a harmful outcome by promoting harmful medical practices or providing incomplete information that could lead to panicked action against the current medical wisdom. The model analyzes the characteristics of people who retweet the article, the pattern of how the article is retweeted, and what twitter users say when retweeting the article.

The study aims to support future work to identify and reduce the spread of panic-inducing misinformation and disinformation in an effort to help authorities better respond to health-threatening epidemics and pandemics.

**Index Terms**—medical informatics, propagation, Twitter, medical advice, categorization, random forest, machine learning, social media, detection

## I. INTRODUCTION

Online social media platforms from the likes of Facebook and Instagram to the likes of Twitter all contain innumerable numbers of fake profiles, troll accounts, and fake news posts, that lead to a rapid spread of misinformation and disinformation. One of the most consequential times for the spread of misinformation and disinformation is in the middle of states of emergencies, and is due to the lack of published, explicit responses from governmental and/or authoritative sources to questions from the public. In these times, some social media users take to publishing their own “armchair expert” opinions on social media while others take to social media to find and solicit answers to the questions they don’t have the information to answer. In these circumstances, the ability for a computer to discern whether an article represents sound advice is important: for one, the social media platform itself can flag posts that appear suspicious in order to curb the spread of misinformation while humans can fact-check the information.

In the realm of biomedicine, the impact of disinformation becomes even more important during times of medical crises. During the 2020 COVID-19 pandemic, the virus was not the only cause of pandemic-related death: misinformation and misinterpretation of legitimate information by officials of all levels led to the death of at least one man, although many more are likely to have been harmed by misinformation [1].

## II. RELATED WORKS

Previous work by Aphinyanaphongs and Aliferis analyzed the validity and trustworthiness of blog posts giving medical

advice for cancer treatments on the internet [2]. This work utilized an analysis of the content of the article or blog post to determine its authority based on the language used. This work is not directly applicable to information exchanged on social media because the content of social media posts on Twitter, for example, must be shorter than 280 (previously 140) characters. This work could be used to analyze the actual article in question, however. Later work by Liu Yang and Yi-Feng Brook Wu during the 2016 US Presidential election season investigated the effectiveness of modeling the propagation via shares of an article for determining whether or not an article was “fake news” [3]. Their findings opened up the possibility for an article or blog post on Twitter to be classified as little as five minutes after being posted based on the characteristics of the people sharing the article as well as the pattern of propagation, an idea also explored by Kwon, Cha, and Jung [4]. The integration of multiple models which categorize not just the temporal forces in propagation but also the characteristics of the people retweeting a given article allowed them to see much higher precision and recall in much less time than previous efforts utilizing just one method or the other.

Despite the apparent dangers of receiving information in emergency situations from social media, Alexander found that social media was used in seven different ways during crises: for “listening to public debate, monitoring situations, extending emergency response and management, crowd-sourcing and collaborative development, creating social cohesion, furthering causes (including charitable donation) and enhancing research” [5]. Because there are a myriad of benefits to social media, in addition to its necessity simply because of prevalence, Alexander postulates that the downsides—especially spreading of rumors and false information—must be endured or dealt with, as “the incorporation of social media into pre-existing emergency management systems is inevitable.” This work shows that there must be additional methods for curtailing rampant spread of misinformation on social media in a way that does not hinder the freedom of expression found on such networks.

## III. HYPOTHESIS

It is this researcher’s goal to test the efficacy of existing methods of detecting fake news and verify whether the existing methodologies show promise in the detection of misinforma-

tion stemming from the user characteristics of the tweets. In this researcher’s opinion, such work could be effectively implemented a few different ways, each as a means to control the spread of misinformation: for one, it could be a user-installed browser extension which could be invoked to determine the reliability of a tweet in question, or it could be implemented in the style of stock market circuit breakers which pause trading on the stock market to prevent panic [6]. Such a system could, when detecting patterns of unhelpful information, alert relevant content moderators and pause sharing on the article while its veracity is ascertained.

#### IV. METHODS

Tweets were collected from *COVID-19: The First Public Coronavirus Twitter Dataset* [7], and were downloaded, cataloged and sorted as a preparatory stage of this project. The initial focus was to find the the most popular Coronavirus-related tweets on Twitter as measured by number of retweets. Next, the articles and source tweets were categorized into three groups: “helpful”, “unhelpful”, and “trash”. Tweets which advocated, promoted, or explained helpful social distancing measures, such as [8], were rated “helpful”. Tweets which advocated bad habits, such as [9], were rated “unhelpful”. This category included tweets such as [10], as in the use case of a browser extension for users, a false negative (resulting in believing harmful advice) is far more impactful than a false positive (where the user would, on the recommendation of the tool, not take the advice of a supposedly sarcastic post). Tweets falling in the third category of “trash” were mostly related to album or movie releases, as these tweets garner a large number of retweets and were included in the COVID-19 dataset due to coronavirus hashtags written in the comments.

The test and training sets were isolated from the top 300 tweets, sorted by number of retweets. 71 tweets were thrown out as trash, while 130 were rated “unhelpful” and 99 were rated “helpful”. The average user characteristics (listed at the end of the next paragraph) for the first 40 retweets were added to each annotated tweet. After this step, the tweets were filtered again to remove any tweets for which the average user characteristics were missing from the dataset (which is caused by the way Twitter’s Search API handles retweets without comments, which meant the users were missing in the original dataset and is not correctable without a professional subscription to Twitter’s API). The final total number of analyzed tweets included 69 marked as helpful and 87 marked as unhelpful, for a grand total of 156 tweets.

A random forest was chosen for the classification, as it showed better performance than a logistic regressive model also tested. 126 tweets were chosen for the training set, with 30 left for the testing set (a 80/20 split). The model was fed the average user characteristic vector (containing 8 features) and a label, either 0 for so-called “helpful” tweets or 1 for an “unhelpful” tweet. Tweaking of the cost complexity pruning alpha of the random forest resulted in a trained model that displayed negligible signs of overfitting when the test and training sets were compared, and further tweaking of the

input user characteristic vector showed that the most pertinent inputs were the length of the user’s description, the number of followers, and the number of statuses posted by the user. Additional input features resulted in either equal or lower quality results. The additional features tested included the length of the user’s screen name, the number of friends, the age of the user’s account, whether or not the account was verified, and whether the user had changed the default profile information. The code used to process the raw tweet data as well as the code to train the model can be found at [11].

#### V. RESULTS

The final trained model displayed a total area under the ROC curve of 0.755 for the test set, which is enough to discount randomness, but not enough to use this model for a business purpose. Using this model for the full set of data results in a similar AUC of 0.720. The ROC curves are plotted for the full set (fig. 1) and the test set (fig. 2). The precision and recall are detailed in table I.

	Precision	Recall	F1-score	Support
0	0.60	0.69	0.64	13
1	0.73	0.65	0.69	17

TABLE I  
PRECISION AND RECALL FOR MODEL ON TEST SET




Fig. 1. Full Set (Train & Test) ROC

#### VI. ANALYSIS AND LIMITATIONS

There are several limitations with the data presented in this study. For one, as was mentioned in the Methods section, the dataset is missing retweets for certain tweets, corresponding to




Fig. 2. Test Set ROC

those which were retweeted without any additional text. This is a major limitation of the study, but can only be overcome with access to the premium twitter API, as the tweets used in the study were too old to use the free retweet lookup functionality in the API.

Additionally, the only characteristics the end model used to predict were the average retweeter's description length, number of followers, and number of statuses posted. These characteristics are likely to bias the results in favor of established and popular users on Twitter, becoming a rough measure of how much "cred" the users who retweet a post have; however, this is not necessarily a bad thing as these accounts will often come from reliable news organizations and world thought leaders who hopefully are more likely to fact-check before rampantly retweeting. Such an assumption would require an entire study on its own to confirm.

Adding additional measures to the model, such as combining the results of the user characteristic measure with those from text categorization of the post or article's content may prove to be a more reliable measure; however, this comes with the difficulty of parsing image and video content as well, as the researcher anecdotes that several of the most helpful posts contained videos of doctors or researchers breaking down complex topics about the pandemic.

Additionally, the quality of this study and model can be thought of more as a representation of how good a random forest can predict whether Matthew Leon would find a given tweet "helpful" or "unhelpful", as the manual classifications contain his personal biases towards certain content. Manual classifications of only one non-expert source will certainly contain biases, especially when much of the content rated con-

tains partisan political statements which have various shades of truth filtered to suit political agendas rather than straight helpful or unhelpful medical advice. A more accurate model for determining the quality of a tweet's medical advice may be made up of several layers: one of which would detect whether or not a tweet contains political content, another which would predict whether a tweet contains advice or a command, and another to determine whether or not that command is in-line with established practice, advises heeding the advice of medical professionals, or spreads panic. Unfortunately, this type of model would likely require training on an established set of crisis-specific data available after the crisis is complete, and may not be useful for emerging situations. In such cases, a more generic model like that used in this study, where the prediction is based upon factors other than the content, may be more helpful.

Finally, this model was not tested in tweets that were unrelated to the COVID-19 crisis, and thus testing such tweets may result in highly erratic predictions. This is an area of further research to investigate the correlation between a tweet's content and the characteristics of users who retweet it. The fact that 20% of the top 300 tweets that were rated being labeled "trash" indicate that automated systems would need to have a countermeasure to deal with this case.

## VII. CONCLUSION

This study set out to show that harmful medical advice could be thwarted by a machine learning model tuned to characteristics about users retweeting a given tweet. While definitive claims applying the results may not be made, the results are indicative of potential for more in-depth analysis factoring more features into account.

## REFERENCES

- [1] "Coronavirus: Man dies taking fish tank cleaner as virus drug," *BBC*, 2020. [Online]. Available: <https://www.bbc.com/news/52012242>
- [2] Y. Aphinyanaphongs and C. Aliferis, "Text categorization models for identifying unproven cancer treatments on the web," *Medinfo 2007: Proceedings of the 12th World Congress on Health (Medical) Informatics; Building Sustainable Health Systems*, p. 968–972, 2007. [Online]. Available: <https://search.informit.com.au/documentSummary;dn=781298524575088;res=IELHEA>
- [3] Y. Liu and Y.-F. Wu, "Early detection of fake news on social media through propagation path classification with recurrent and convolutional networks," in *AAAI*, 12 2018. [Online]. Available: <https://www.aaai.org/ocs/index.php/AAAI/AAAI18/paper/view/16826>
- [4] S. Kwon, M. Cha, and J. Kyomin, "Rumor detection over varying time windows," *PLOS ONE*, vol. 12, no. 1, pp. 1–19, 01 2017. [Online]. Available: <https://doi.org/10.1371/journal.pone.0168344>
- [5] D. E. Alexander, "Social media in disaster risk reduction and crisis management," *Science and Engineering Ethics*, vol. 20, no. 3, pp. 717–733, Sep. 2014. [Online]. Available: <https://doi.org/10.1007/s11948-013-9502-z>
- [6] A. Schneider and S. Horsley. (2020). [Online]. Available: <https://www.npr.org/2020/03/09/813682567/how-stock-market-circuit-breakers-work>
- [7] E. Chen, K. Lerman, and E. Ferrara, "COVID-19: The first public coronavirus twitter dataset," 2020. [Online]. Available: <https://arxiv.org/abs/2003.07372v1>

- [8] B. Obama. (2020) If you're wondering whether it's an overreaction to cancel large gatherings and public events (and I love basketball), here's a useful primer as to why these measures can slow the spread of the virus and save lives. we have to look out for each other. [Online]. Available: <https://twitter.com/BarackObama/status/1238238576141352966>
- [9] T. Jardin. (2020) Coronavirus is everywhere. BOOK THAT FLIGHT. take that trip. you probably won't die from it, but even if you do... you wanna die with the eiffel tower in the background, or you wanna die with I-20 in the background? you decide. [Online]. Available: <https://twitter.com/JardinTaylor/status/1237429345632882688>
- [10] K. Kinder. (2020) ALERT!!!!!! the corona virus can be spread through money. if you have any money at home, put on some gloves, put all the money in to a plastic bag and put it outside the front door tonight. I'm collecting all the plastic bags tonight for safety. think of your health. [Online]. Available: <https://twitter.com/stayfrea/status/1235256530728972290>
- [11] M. Leon. (2020) Supplementary materials for "detecting harmful medical advice by analyzing the characteristics of retweeters". [Online]. Available: <https://github.com/leonm1/medical-advice-twitter-project>